

Vanna Winland

vgwinland@gmail.com | (330) 703-2490 | Akron, OH | [linkedin.com/in/vwinland](https://www.linkedin.com/in/vwinland) | vwinland.github.io

SUMMARY

Developer Relations and content strategy professional who translates complex technology for any audience — technical writers, PMMs and content strategists, or businesses figuring out how to enable their teams with AI — backed by independent SEO/AEO research that puts published work on page one of Google for competitive terms. Founded a DevRel function at an early-stage graph database and AI infrastructure startup, and now shapes AI and developer content strategy at IBM, building the team's documentation and contribution infrastructure while shipping working software like RAG applications and multi-agent systems, not just writing about them. Research like a journalist, build like an engineer, diving into unfamiliar codebases to find answers rather than settle for secondhand documentation. Independently writes and publishes essays on AI and software engineering, syndicated to HackerNoon, favoring grounded, practitioner-focused analysis over industry hype.

EXPERIENCE

Developer Advocate & Technical Content Creator | IBM2024 – Present

- Write and publish technical content on AI agents, LLM orchestration, and generative AI systems for ibm.com/think, translating complex implementations for developer and technical business audiences.
- Partner with SEO stakeholders and conduct independent keyword and topic research using Ahrefs and Clearscope to validate traffic potential before writing, optimizing for both traditional SEO and answer engine optimization (AEO); majority of published articles rank on page one of Google for target terms, including competitive terms like “LLM orchestration.”
- Research and evaluate emerging AI frameworks and infrastructure firsthand — including CrewAI, agent communication standards (A2A, ACP), IBM watsonx.ai and watsonx Orchestrate, and IBM's agent telemetry and observability tooling — to ground technical coverage in direct, hands-on use rather than secondhand documentation.
- Cover real-time event-streaming and data-streaming platforms, investigating managed vs. self-managed infrastructure (e.g., Confluent/Kafka) to ground reporting on distributed data systems in direct, hands-on use.
- Independently wrote the script for a video walkthrough on agentic architecture, partnering with IBM's A/V team on recording.
- Write tutorials for IBM Bob, covering project-level rules and configuration for AI coding agents.
- Maintain a documentation site and GitHub repository as primary content infrastructure — implementing security and contribution safeguards, and designing a new authoring and contribution workflow (pending manager review) — while coordinating contributions and keeping published work accurate as the technology evolves.
- Collaborate with engineering and product teams to pressure-test technical accuracy before publication.

Founding Developer Relations Engineer | GraphGrid / AtomRain2020 – 2024

- Built the company's content and documentation strategy from scratch as the first dedicated DevRel hire — designing the documentation site's information architecture and maintaining code samples across all product versions, and establishing editorial templates and publishing workflows for an open-source graph database product and its commercial platform.
- Set language standards for consistency across product and marketing in collaboration with the marketing team, including developing an internal terminology glossary to enforce legal and accuracy requirements (e.g., restricted product/legacy-name usage) across all published content.
- Reverse-engineered product behavior by reading directly through the Java codebase without prior Java experience, using that research to accurately document how the API and product actually worked.
- Wrote API documentation for 12 core product modules, including example requests and responses.
- Built and maintained curated Postman collections for distinct use cases, including handoff packages for penetration testing and customer/developer onboarding.
- Produced long-form technical guides and documentation spanning APIs, UI dashboards, Spring Cloud/AWS configuration services, Kafka- and RabbitMQ-based change data capture and messaging, NLP data extraction and graph data, Airflow-orchestrated ML training, Spark-based machine learning, and Elasticsearch search indexing — learning each platform hands-on before writing.
- Created a video content series — writing scripts, narrating walkthroughs, and recording using Adobe Creative Cloud and OBS — explaining complex platform concepts to mixed technical audiences.
- Reviewed and edited whitepapers and social media content submitted by contracted technical writers for accuracy.
- Initiated social/developer forums and product presence on Reddit, Discord, Slack, and Stack Overflow, building direct relationships with an open-source developer community.
- Organized curated client demonstrations and conducted market research to identify prospects and industry trends.
- Served as point of contact for technical training and troubleshooting; mentored student interns through onboarding.
- Contributed to product roadmap and feature scoping in collaboration with the product team.

PROJECTS

Ask My Archive | Python, Chroma, Sentence-Transformers, Streamlit, Anthropic API, Hugging Face Inference API

[Source](#) | [Live App](#)

- Built and deployed an end-to-end RAG (retrieval-augmented generation) application that answers questions against my own published writing, grounding every response in retrieved excerpts with citations instead of model-generated guesses.
- Designed a metadata schema resolving syndicated essays published under different titles across four platforms (personal blog, HackerNoon, Medium, IBM Think) into single canonical entries, and built a persistent Chroma vector store indexing 43 documents (1,749 chunks).
- Built a swappable generation backend (Claude API, local Ollama, Hugging Face Serverless Inference) sharing one retrieval and citation-parsing pipeline; deployed a public Streamlit chat interface and debugged production issues across two hosting platforms.

A2A Protocol Agent Communication System | Python, BeeAI, A2A Protocol, Ollama

[Source](#)

- Built a client-server AI agent system for the IBM Think tutorial “Use the A2A protocol for AI agent communication,” implementing an A2A-compliant server that exposes a tool-using BeeAI agent (web search, weather, Wikipedia, step-by-step reasoning) over HTTP with a discoverable AgentCard endpoint.
- Built a separate async terminal client exchanging JSON-RPC/SSE messages with the server, demonstrating decoupled agent-to-agent communication running against a locally hosted Ollama model.

ChatDev on IBM watsonx.ai | Python, CAMEL Framework, IBM watsonx.ai, Hugging Face

[Source](#)

- Adapted the open-source ChatDev multi-agent framework to run on IBM watsonx.ai for the IBM Think tutorial “Use ChatDev ChatChain for agent communication on IBM watsonx.ai,” building a custom model backend wrapper and modifying prompt templates, model selection logic, and token-limit configuration to integrate Llama-4-Maverick via watsonx's Hugging Face integration.
- Enabled a five-agent virtual software team (CEO, CTO, Designer, Tester, Programmer) to design, code, and test a working app end to end from a single prompt, porting the framework using an AI-assisted, prompt-driven build approach.

SKILLS

Content strategy & editorial systems: documentation architecture, style/language standards, SEO/AEO, content governance
 AI/ML & agents: RAG (Chroma, Sentence-Transformers, retrieval + citation pipelines), agent frameworks (BeeAI, CrewAI, CAMEL/ChatDev), agent communication protocols (A2A, ACP), LLM orchestration, agent telemetry & observability, watsonx.ai

Data & infrastructure: distributed messaging & streaming (Kafka, RabbitMQ, Confluent), cloud infrastructure (AWS), Spark, Elasticsearch, Airflow

Languages & tools: Python, GitHub (incl. repo security/contribution controls), Streamlit, Hugging Face, Anthropic API, Ollama, Ahrefs, Clearscope, Adobe Creative Cloud, OBS, Postman

Leadership & operations: cross-team coordination, workflow design, mentoring, standup facilitation, open-source community building

EDUCATION

B.A., English (Literature) — University of Akron

Full Stack Web Development Certificate — Flatiron School

WRITING & PUBLICATIONS

[IBM Think Author Profile](#) | [Personal Blog](#)

- “Industrialization of Software Engineering” — an ongoing independent essay series examining how AI is reshaping software development, spanning code review, governance, requirements engineering, and enterprise migration patterns.
- Published 18 essays on HackerNoon in under three months (June–August 2026), generating 20,800+ combined page views, with individual pieces reaching as many as 8,800+ reads.
- Four essays selected as HackerNoon Top Stories — homepage placement plus syndication across Twitter, Noonification, Facebook, LinkedIn, and Instagram, in 12 translations: “A24's AI Backlash Is an Old Creative Argument in New Clothes” (also cited by HackerNoon's CEO in a public roundup), “Don't Buy an Uncensored AI on a Flash Drive,” “The Problem With Using AI to Review AI-Written Code” (8,810 page views), and “I Built a RAG Tool Over My Own Writing—Here's What I Learned.”
- “What Is an Agent, Actually?” was the subject line of HackerNoon's subscriber newsletter for the edition featuring it.
- Four essays trended in HackerNoon's Daily Tech Beat for multiple consecutive days: the AI-backlash essay, “Don't Buy an Uncensored AI on a Flash Drive,” “The IDE Is No Longer Just a Place to Write Code,” and “What Is an Agent, Actually?”
- “The Problem With Autogenerated Docs” drew direct professional outreach from a senior technical publications manager at Oracle.
- Published essays and technical articles consistently rank on page one of Google search results for their target terms.