

Vanna Winland

vgwinland@gmail.com | (330) 703-2490 | Akron, OH | <https://www.linkedin.com/in/vwinland/> | vwinland.github.io

SUMMARY

Developer Relations and content strategy professional who designs documentation systems and internal workflows from the ground up. Founded a DevRel function at an early-stage graph database and AI infrastructure startup, and now shapes AI and developer content strategy at IBM. Research like a journalist, build like an engineer, with a technical fearlessness that means diving into unfamiliar codebases to find answers rather than settling for secondhand documentation. Independently writes and publishes essays on AI and software engineering, syndicated to HackerNoon, that favor grounded, practitioner-focused analysis over industry hype.

EXPERIENCE

Developer Advocate & Technical Content Creator | IBM

2024 – Present

- Write and publish technical content on AI agents, LLM orchestration, and generative AI systems for ibm.com/think, translating complex implementations for developer and technical business audiences.
- Partner with SEO stakeholders and conduct independent keyword and topic research using Ahrefs and Clearscope to validate traffic potential before writing; majority of published articles rank on page one of Google for target terms, including competitive terms like “LLM orchestration.”
- Optimize content for both traditional SEO and answer engine optimization (AEO).
- Research and evaluate emerging AI frameworks firsthand, including early-stage projects like CrewAI, building working examples to ground coverage in real systems.
- Cover IBM watsonx.ai, [watsonx Orchestrate](https://watsonx.com/orchestrate), open-source LLM tooling, and enterprise data workflows, developing platform fluency through direct use.
- Write tutorials for IBM Bob, covering project-level rules and configuration for AI coding agents.
- Maintain a documentation site and GitHub repository as primary content infrastructure, coordinating contributions and keeping published work accurate as the technology evolves.
- Implemented security controls and contribution safeguards on a public open-source GitHub repository.
- Designed and documented a new authoring and contribution workflow for the team's content and repository processes; currently pending manager review before rollout.
- Collaborate with engineering and product teams to pressure-test technical accuracy before publication.

Founding Developer Relations Engineer | GraphGrid / AtomRain

2020 – 2024

- Built the company's content and documentation strategy from scratch as the first dedicated DevRel hire, establishing editorial architecture, templates, and publishing workflows.
- Ideated, developed, and maintained the technical documentation site and code samples across all product versions and public/internal forms, designing the site's information architecture from the ground up.
- Set language standards for consistency across product and marketing in collaboration with the marketing team, including developing an internal terminology glossary to enforce legal and accuracy requirements (e.g., restricted product/legacy-name usage) across all published content.
- Reverse-engineered product behavior by reading directly through the Java codebase without prior Java experience, using that research to accurately document how the API and product actually worked.
- Wrote API documentation for 12 core product modules, including example requests and responses.
- Built and maintained curated Postman collections for distinct use cases, including handoff packages for penetration testing and customer/developer onboarding.
- Produced long-form technical guides on APIs, cloud data infrastructure, and graph-based AI systems, learning each platform hands-on before writing.
- Wrote documentation spanning UI dashboards, Spring Cloud/AWS configuration services, Kafka and RabbitMQ-based change data capture, NLP data extraction and graph data, Airflow-orchestrated ML training, Spark-based machine learning, and Elasticsearch search indexing.

- Created a video content series — writing scripts, narrating walkthroughs, and recording using Adobe Creative Cloud and OBS — explaining complex platform concepts to mixed technical audiences.
- Reviewed and edited whitepapers and social media content submitted by contracted technical writers for accuracy.
- Initiated social/developer forums and product presence on Reddit, Discord, Slack, and Stack Overflow.
- Organized curated client demonstrations and conducted market research to identify prospects and industry trends.
- Served as point of contact for technical training and troubleshooting; mentored student interns through onboarding.
- Contributed to product roadmap and feature scoping in collaboration with the product team.

SKILLS

- Content strategy & editorial systems: documentation architecture, style/language standards, SEO/AEO, content governance
- Technical range: LLM orchestration, AI agent frameworks, watsonx.ai, cloud infrastructure (AWS), Kafka/RabbitMQ, Spark, Elasticsearch, Airflow
- Tools: GitHub (incl. repo security/contribution controls), Ahrefs, Clearscope, Adobe Creative Cloud, OBS, Postman
- Leadership & operations: cross-team coordination, workflow design, mentoring, startup facilitation

EDUCATION

- B.A., English (Literature) — University of Akron (2013-2017)
- Flatiron School – Full Stack Web Development (2019-2020)

WRITING & PUBLICATIONS

- “Industrialization of Software Engineering” — an ongoing independent essay series examining how AI is reshaping software development, spanning code review, governance, requirements engineering, and enterprise migration patterns.
- Published 13 essays in under two months (June–July 2026), generating 6,300+ combined page views, with individual pieces reaching as many as 1,769 reads.
- Two essays selected as HackerNoon Top Stories: the essay on AI backlash in creative industries (also cited by HackerNoon's CEO in a public roundup) and “Don't Buy an Uncensored AI on a Flash Drive.” Both Top Story selections are featured in HackerNoon's subscriber newsletter.
- “What Is an Agent, Actually?” was the subject line of HackerNoon's subscriber newsletter for the edition featuring it.
- Four essays trended in HackerNoon's Daily Tech Beat for multiple consecutive days: the AI-backlash essay, “Don't Buy an Uncensored AI on a Flash Drive,” “The IDE Is No Longer Just a Place to Write Code,” and “What Is an Agent, Actually?”
- “The Problem With Autogenerated Docs” drew direct professional outreach from a senior technical publications manager at Oracle.
- Published essays and technical articles consistently rank on page one of Google search results for their target terms.